

Additional file 1

Assay resolution governs the transfer of molecular dependence

Anonymous Authors

Sections 1–6 document the original fixed-margin benchmark, including its estimators, study protocols, and all recorded outcomes. Section 7 gives the theory and retrospective analyses of assay resolution and recipient coexpression reported in the current manuscript.

Contents

1	Coupling coordinates	2
1.1	Definition	2
2	Binary fixed-margin estimator	2
2.1	Binary scaling identity	2
2.2	Fixed-margin likelihood	3
2.3	Hierarchical transfer objective	3
2.4	Source-only graphs and tuning grids	4
2.5	Transfer risk and recipient drift	4
2.6	Decision rule for table prediction	5
2.7	Matched signed-statistic transfer	6
2.8	Fitted interaction baselines	6
2.9	Conditional random-effects prediction	7
2.10	Ground-truth estimator simulation	7
2.11	Recipient heterogeneity	7
2.12	Transfer-risk stress test	7
3	Secondary multistate methods	8
3.1	Permutation-centered fields	8
3.2	Structured field estimator	9
4	Public binary studies	9
4.1	COMBAT and BMBC development panels	10
4.2	Stephenson held-site confirmation	10
4.3	Matched predictive reanalysis	10
4.4	GSE239452 cohort replication	11
4.5	GSE314416 held-pool attempt	11
4.6	GSE317605 longitudinal calibration	12
5	Binary conditional transfer results	12
5.1	Held outcomes	12
5.2	Stephenson linkage and lineage sensitivity	12
5.3	Composition-preserving and protein-threshold sensitivities	13
5.4	Selected-panel signal sensitivity	13
5.5	Fitted interaction baselines	13
5.6	Matched development comparisons	13
5.7	Public analysis sequence	14
6	Secondary multistate results	14
6.1	Data and preprocessing	14
6.2	Predictive comparisons	15
6.3	PerturbSci sensitivity and functional analysis	16

7 Assay resolution and recipient coexpression	16
7.1 Sharp coarse prediction bound	16
7.2 Response to a transferred interaction	17
7.2.1 Query discrepancy from covariance mismatch	18
7.3 Matched resolution and covariance analyses	18
7.4 Recipient-specific and pooled coexpression	19
7.5 Query ranking and uncertainty	19
7.6 Relation to full-profile interaction fitting	19
7.7 Computational verification	19
8 Limitations	20
9 Implementation and reproducibility	20

1 Coupling coordinates

1.1 Definition

Let $P \in (0, 1)^{r \times s}$ be a joint probability table, and let $B_r \in \mathbb{R}^{r \times (r-1)}$ and $B_s \in \mathbb{R}^{s \times (s-1)}$ have orthonormal columns spanning the complements of the all-ones vectors. The population coupling field is

$$\Theta(P) = B_r^\top \log(P) B_s. \quad (1)$$

We use Helmert bases. Other orthonormal contrast bases rotate the $(r-1)(s-1)$ coordinates without changing Frobenius norms or pairwise Frobenius distances.

Proposition S1 (Marginal invariance and independence). *For positive vectors $a \in \mathbb{R}^r$ and $b \in \mathbb{R}^s$, define*

$$P' = \frac{\text{diag}(a) P \text{diag}(b)}{\sum_{u,v} a_u P_{uv} b_v}.$$

Then $\Theta(P') = \Theta(P)$. Moreover, $\Theta(P) = 0$ if and only if P is an independent joint law.

Proof. Writing c for the normalizing constant gives

$$\log P' = \log P + (\log a) \mathbf{1}_s^\top + \mathbf{1}_r (\log b)^\top - \log(c) \mathbf{1}_r \mathbf{1}_s^\top.$$

The added terms are annihilated on the left by $B_r^\top$ or on the right by B_s , proving invariance. Factorization $P_{uv} = p_u q_v$ makes $\log P$ a sum of row and column terms, so $\Theta(P) = 0$. Conversely, $\Theta(P) = 0$ removes the interaction component in the orthogonal decomposition into row and column terms. Hence $\log P_{uv} = \alpha_u + \beta_v$, and P factorizes after normalization. $\square$

2 Binary fixed-margin estimator

The original binary benchmark estimator conditions RNA-protein tables on their observed margins, eliminating nuisance marginal parameters. Predictions use the recipient margins without a pseudocount.

2.1 Binary scaling identity

For

$$P = \begin{pmatrix} p_{00} & p_{01} \\ p_{10} & p_{11} \end{pmatrix}, \quad B_2 = 2^{-1/2} (1, -1)^\top,$$

Eq. (1) is scalar and satisfies

$$\Theta(P) = \frac{1}{2} \log \frac{p_{00} p_{11}}{p_{01} p_{10}}. \quad (2)$$

The conventional log odds ratio is therefore $\theta = 2\Theta(P)$. Estimation uses θ ; reported field coordinates use the Helmert scale in Eq. (2).

2.2 Fixed-margin likelihood

For donor d and ordered RNA-protein pair e , write the 2×2 table as

$$N_{de} = \begin{pmatrix} A_{de} & r_{0,de} - A_{de} \\ c_{0,de} - A_{de} & n_{de} - r_{0,de} - c_{0,de} + A_{de} \end{pmatrix},$$

where $r_{0,de}$ and $c_{0,de}$ are the first row and column totals and $c_{1,de} = n_{de} - c_{0,de}$. Conditional on these margins,

$$\Pr_{\theta_{de}}(A = a \mid r_0, c_0, n) = \frac{\binom{c_0}{a} \binom{c_1}{r_0 - a} \exp(\theta_{de} a)}{Z_{de}(\theta_{de})}, \quad Z_{de}(\theta) = \sum_{a \in \mathcal{A}_{de}} \binom{c_0}{a} \binom{c_1}{r_0 - a} \exp(\theta a), \quad (3)$$

on the feasible integer support $\mathcal{A}_{de}$. Omitting a constant independent of θ , the negative log likelihood is

$$\ell_{de}(\theta) = \log Z_{de}(\theta) - A_{de}\theta, \quad \ell'_{de}(\theta) = \mathbb{E}_{\theta}[A] - A_{de}, \quad \ell''_{de}(\theta) = \text{Var}_{\theta}(A). \quad (4)$$

An informative table has at least two values in $\mathcal{A}_{de}$, making $\ell''_{de} > 0$ and the conditional odds ratio identifiable.

Hierarchical conditional and random-effects log-linear models share odds-ratio information across 2×2 strata [1, 2, 3]. We estimate donor-specific and population log odds jointly across the entity panel, then predict recipient tables at their observed margins.

2.3 Hierarchical transfer objective

Let D and E denote the numbers of source donors and ordered marker pairs, θ_{de} a donor-specific log odds ratio, and μ_e the transferable population log odds. Two marker hypergraphs induce normalized Laplacians L_R and L_A for the RNA and protein axes. Their product-space penalty is

$$L = L_R \otimes I + I \otimes L_A. \quad (5)$$

The estimator minimizes the penalized composite conditional objective

$$Q(\theta, \mu) = \frac{1}{D} \sum_{d=1}^D \sum_{e=1}^E \ell_{de}(\theta_{de}) + \frac{\gamma}{2} \sum_{d,e} (\theta_{de} - \mu_e)^2 + \frac{\rho}{2} \|\mu\|_2^2 + \frac{\tau}{2} \mu^\top L \mu. \quad (6)$$

The likelihood sum is composite because entity pairs within a donor share cells and markers. We express the penalty grid through dimensionless inputs

$$\gamma = \eta s_\eta / D, \quad \rho = \lambda_0 s_\mu, \quad \tau = \lambda_G s_\mu,$$

where s_η is the median null Fisher information across informative donor-entity tables and s_μ is the median information after donor averaging. This scaling makes the grid invariant to duplication of the complete donor panel.

Its gradients are

$$\frac{\partial Q}{\partial \theta_{de}} = \frac{\mathbb{E}_{\theta_{de}}[A] - A_{de}}{D} + \gamma(\theta_{de} - \mu_e), \quad (7)$$

$$\nabla_\mu Q = -\gamma \sum_d (\theta_d - \mu) + \rho \mu + \tau L \mu. \quad (8)$$

The Hessian has diagonal θ block $\text{diag}\{\text{Var}_{\theta_{de}}(A)/D + \gamma\}$, θ - μ entries $-\gamma$, and μ block $D\gamma I + \rho I + \tau L$. Its quadratic form is

$$\sum_{d,e} \frac{\text{Var}_{\theta_{de}}(A)}{D} u_{de}^2 + \gamma \sum_{d,e} (u_{de} - v_e)^2 + \rho \|v\|_2^2 + \tau v^\top L v. \quad (9)$$

Proposition S2 (Finite identifiable fit). *If $\gamma > 0$, $\rho > 0$, and $\tau \geq 0$, Eq. (6) has a unique finite minimizer. For every informative donor-entity table, its unpenalized conditional parameter is identifiable.*

Proof. Each likelihood term is convex by Eq. (4), and the product Laplacian is positive semidefinite. Equality to zero in Eq. (9) requires $u_{de} = v_e$ through the γ term and $v = 0$ through the ρ term, hence $u = 0$. The objective is strictly convex and coercive. The final statement follows from $\ell''_{de} > 0$ on every informative support. $\square$

Newton updates eliminate donor-specific parameters by the Schur complement. The population diagonal then receives

$$\sum_d \frac{\gamma \{\text{Var}_{\theta_{de}}(A)/D\}}{\gamma + \text{Var}_{\theta_{de}}(A)/D}$$

in addition to $\rho I + \tau L$. Backtracking enforces descent, and fits that exceed prespecified gradient-norm or condition-number thresholds are rejected.

For a recipient donor, Eq. (3) with log odds $\alpha \hat{\mu}_e$ and the recipient margins yields a conditional-expectation prediction that preserves both margins. Development donors determine α and all penalties.

2.4 Source-only graphs and tuning grids

Graphs used only samples available at the corresponding selection stage. Stephenson and COMBAT represented markers by RNA detection prevalence and mean log-transformed ADT panel fractions across source sample-by-cell-type strata. GSE239452 used the corresponding donor-level profiles. Standardized profiles defined Euclidean k -nearest-neighbor graphs; the union of directed neighbor relations formed undirected two-vertex hyperedges. BMMC used prevalence across source donor-by-lineage strata, with each marker and its k nearest neighbors forming one unweighted hyperedge. The resulting normalized RNA and ADT hypergraph Laplacians entered Eq. (5). Cross-validation rebuilt both graphs within each calibration fold, and the final fit used all permitted development samples.

Table 1 gives the complete binary grids. Redundant values of k were omitted when $\lambda_G = 0$. The statistic comparator crossed Pearson and deviance statistics, raw and conditional-null-centered forms, and the same transport grid; GSE239452 evaluated raw statistics only.

Table 1: Binary estimator grids used for the reported analyses.

Analysis	Neighbors k	Heterogeneity η	Ridge λ_0	Graph λ_G	Transport α
Stephenson	1, 2	0.1, 1, 10	0.01, 0.1	0, 0.1, 1	0.5, 0.75, 1, 1.25
GSE239452	1, 2	0.1, 1, 10	0.01, 0.1	0, 0.1, 0.3, 1	0.5, 0.75, 1, 1.25
COMBAT	1, 2	0.1, 1, 10	0.01, 0.1	0, 0.1, 0.3, 1	0.5, 0.75, 1, 1.25
development					
BMMC	1, 2, 3	0.1, 1, 10	0.01, 0.1	0, 0.1, 0.3, 1	0.5, 0.75, 1, 1.25
development					

2.5 Transfer risk and recipient drift

To separate source estimation error from recipient variation, fix informative margins $M = (r_0, r_1, c_0, c_1)$ with total N and endpoints

$$L_M = \max(0, r_0 + c_0 - N), \quad U_M = \min(r_0, c_0)$$

for the feasible upper-left count. Write Eq. (3) as

$$p_\theta(a | M) = h_M(a) \exp\{\theta a - \psi_M(\theta)\}, \quad I_M(\theta) = \psi_M''(\theta) = \text{Var}_\theta(A | M).$$

Proposition S3 (Fixed-margin transfer risk). *Let $\mathcal{S}$ contain all source information, and suppose a recipient log odds satisfies*

$$\theta_q = \mu + \xi_q, \quad \mathbb{E}(\xi_q | \mathcal{S}, M) = 0, \quad \mathbb{E}(\xi_q^2 | \mathcal{S}, M) = \sigma_M^2.$$

The source estimate $\hat{\mu}$ and multiplier α are $\mathcal{S}$ -measurable, and the recipient prediction uses $\hat{\theta} = \alpha\hat{\mu}$. Assume that θ_q and $\hat{\theta}$ lie almost surely in a compact interval $\mathcal{T}$. Define

$$\underline{I}_M = \inf_{\theta \in \mathcal{T}} I_M(\theta), \quad \bar{I}_M = \sup_{\theta \in \mathcal{T}} I_M(\theta).$$

Then the conditional excess log loss

$$\mathcal{R}_M = \mathbb{E}[\text{KL}\{p_{\theta_q}(\cdot | M) \| p_{\hat{\theta}}(\cdot | M)\} | \mathcal{S}, M]$$

satisfies

$$\frac{\underline{I}_M}{2} \{(\alpha\hat{\mu} - \mu)^2 + \sigma_M^2\} \leq \mathcal{R}_M \leq \frac{\bar{I}_M}{2} \{(\alpha\hat{\mu} - \mu)^2 + \sigma_M^2\}. \quad (10)$$

Moreover, $I_M(\theta) \leq (U_M - L_M)^2/4$ for every finite θ , so

$$\mathcal{R}_M \leq \frac{(U_M - L_M)^2}{8} \{(\alpha\hat{\mu} - \mu)^2 + \sigma_M^2\}. \quad (11)$$

Divide these bounds by N for excess log loss per cell.

Proof. For canonical parameters θ and η ,

$$\text{KL}(p_\theta \| p_\eta) = \psi_M(\eta) - \psi_M(\theta) - \psi_M'(\theta)(\eta - \theta).$$

Taylor's theorem with integral remainder gives

$$\text{KL}(p_\theta \| p_\eta) = (\eta - \theta)^2 \int_0^1 (1-t) I_M\{\theta + t(\eta - \theta)\} dt.$$

Bounding the information along the segment from θ to η , substituting $\theta = \mu + \xi_q$ and $\eta = \alpha\hat{\mu}$, and taking the conditional expectation prove Eq. (10). Popoviciu's inequality gives $I_M(\theta) = \text{Var}_\theta(A \mid M) \leq (U_M - L_M)^2/4$, which proves Eq. (11). For informative margins, I_M is continuous and positive on the compact interval $\mathcal{T}$, hence $I_M > 0$. $\square$

If α is fixed across source repetitions and $b = \mathbb{E}(\hat{\mu}) - \mu$, the reducible term in Eq. (10) decomposes as

$$\mathbb{E}(\alpha\hat{\mu} - \mu)^2 = \alpha^2 \text{Var}(\hat{\mu}) + \{\alpha b + (\alpha - 1)\mu\}^2. \quad (12)$$

As source estimation becomes consistent and α approaches one, σ_M^2 remains in the conditional log loss of a single transported parameter relative to the donor-specific conditional law. This comparison does not imply a positive lower bound on mean-table prediction regret.

2.6 Decision rule for table prediction

The conditional model defines a distribution over integer tables with the recipient margins. The original binary benchmark predicts its mean table.

Proposition S4 (Conditional mean minimizes expected deviance). *Let N be a random contingency table with fixed margins M and total $n > 0$, and let $\bar{N} = \mathbb{E}(N \mid M)$. For any nonnegative table Q with margins M and $Q_{ij} > 0$ whenever $\bar{N}_{ij} > 0$,*

$$\mathbb{E}\{D(N, Q) \mid M\} - \mathbb{E}\{D(N, \bar{N}) \mid M\} = \frac{2}{n} \sum_{i,j: \bar{N}_{ij} > 0} \bar{N}_{ij} \log \frac{\bar{N}_{ij}}{Q_{ij}} \geq 0. \quad (13)$$

The conditional mean is the unique minimizer on cells with positive conditional mean.

Proof. For any admissible Q ,

$$\mathbb{E}\{D(N, Q) \mid M\} = \frac{2}{n} \mathbb{E} \left(\sum_{i,j: N_{ij} > 0} N_{ij} \log N_{ij} \mid M \right) - \frac{2}{n} \sum_{i,j} \bar{N}_{ij} \log Q_{ij}.$$

The first term is independent of Q . Subtracting the same expression evaluated at $\bar{N}$ gives Eq. (13). Because Q/n and $\bar{N}/n$ are probability tables, the right side is twice their Kullback–Leibler divergence. It is nonnegative and vanishes only when $Q = \bar{N}$ on the support of $\bar{N}$. Every table in the conditional support has margins M , so linearity of expectation gives the same margins to $\bar{N}$. $\square$

For a predictive interaction law G , linearity gives the mean table $Q_G = \int \mathbb{E}_\theta(N \mid M) dG(\theta)$. Proposition S4 applies to this mixture as well as to a single interaction. The following consequence connects error in the interaction distribution to the table loss.

Corollary S1 (Count-uniform stability of table prediction). *Fix binary margins M with total $n > 0$ and feasible upper-left counts $\{L, \dots, U\}$; write $w = U - L$. For probability laws G, H on real log odds with finite second moments, define*

$$P_G(a) = \int p_\theta(a \mid M) dG(\theta), \quad Q_G = \mathbb{E}_{P_G}(N),$$

and define P_H, Q_H analogously. Then

$$\mathbb{E}_{P_G}\{D(N, Q_H) - D(N, Q_G)\} = 2 \text{KL}(Q_G/n \parallel Q_H/n) \leq \frac{w}{2n} W_2^2(G, H) \leq \frac{1}{4} W_2^2(G, H), \quad (14)$$

where W_2 is the quadratic Wasserstein distance. Structural zero cells contribute zero. If $w = 0$, the excess deviance and margin-specific bound vanish. For normal laws, including zero variances, $W_2^2(G, H) = (\mu_G - \mu_H)^2 + (\sigma_G - \sigma_H)^2$. The universal constant $1/4$ cannot be decreased over all binary margins and interaction laws. The laws may depend on the fixed margins M .

Proof. Proposition S4 gives the equality. The noncentral-hypergeometric factorization of Kou and Ying [4], Lemma 2.1 and Theorem 2.1, gives margin-dependent constants $\lambda_j > 0$ such that

$$A - L \stackrel{d}{=} \sum_{j=1}^w B_j, \quad B_j \text{ independent}, \quad p_j(\theta) := \Pr(B_j = 1) = \frac{e^\theta}{e^\theta + \lambda_j}.$$

Their odds parameter is e^θ in our notation. The L deterministic successes correspond to zero roots of the generating polynomial. Define

$$R = \begin{pmatrix} L & r_0 - U \\ c_0 - U & n - r_0 - c_0 + L \end{pmatrix}, \quad C(p) = \frac{1}{2} \begin{pmatrix} p & 1 - p \\ 1 - p & p \end{pmatrix}.$$

The matrix R is nonnegative, has total $n - 2w$, and is independent of θ . The normalized conditional mean therefore decomposes as

$$q_\theta := \frac{\mathbb{E}_\theta N}{n} = \frac{R}{n} + \frac{2}{n} \sum_{j=1}^w C\{p_j(\theta)\}.$$

The first term is omitted when $n = 2w$. Each component has the same KL divergence as its Bernoulli variable:

$$\text{KL}\{C(p)\|C(q)\} = \text{KL}\{\text{Ber}(p)\|\text{Ber}(q)\}.$$

Joint convexity therefore gives

$$\text{KL}(q_\theta\|q_\eta) \leq \frac{2}{n} \sum_{j=1}^w \text{KL}\{\text{Ber}(p_j(\theta))\|\text{Ber}(p_j(\eta))\} \leq \frac{w}{4n} (\theta - \eta)^2.$$

The last inequality follows because Bernoulli log-partition curvature is at most $1/4$ and each component's logit changes by $\theta - \eta$. For any coupling π of G, H , another application of joint convexity gives

$$2 \text{KL}(Q_G/n\|Q_H/n) \leq 2 \int \text{KL}(q_\theta\|q_\eta) d\pi \leq \frac{w}{2n} \int (\theta - \eta)^2 d\pi.$$

Taking the infimum proves the first bound; $2w \leq n$ proves the second. For sharpness, take $n = 2$, $r_0 = c_0 = 1$, $G = \delta_0$, and $H = \delta_h$. Then $W_2^2(G, H) = h^2$ and $2 \text{KL}(Q_G/n\|Q_H/n) = 2 \log \cosh(h/2) = h^2/4 + o(h^2)$. The normal-law expression follows from monotone quantile coupling. $\square$

Different interaction laws can yield the same mean table, so positive heterogeneity need not produce positive table regret.

The bound concerns normalized table prediction, not the full predictive count distribution. The elementary data-processing bound $2 \text{KL}(Q_G/n\|Q_H/n) \leq 2 \text{KL}(P_G\|P_H)$ also holds, but does not give the factor $w/(2n)$ in Eq. (14). For point interaction laws, sufficiency of $\sum_j B_j$ gives the stronger contraction $\text{KL}(q_\theta\|q_\eta) \leq (2/n) \text{KL}(p_\theta\|p_\eta)$. This factor does not extend to arbitrary mixtures: at balanced margins, mixtures concentrated near the two support endpoints can retain almost the same KL divergence in the mean table as in the count law.

Direct summation checked Eq. (14) for 311,328 ordered point-law comparisons across every informative margin with $2 \leq n \leq 24$ and log odds in $\{-12, -6, -2, -0.2, 0, 0.2, 2, 6, 12\}$. A further 1,200 comparisons used seven-atom interaction laws at $n \in \{2, 8, 32, 128, 512, 2048\}$. None exceeded the bound at tolerance 10^{-10} ; the point-law ratio reached 0.99834. The sharpness example approached 0.25 after division by h^2 . These numerical checks use no biological outcomes.

2.7 Matched signed-statistic transfer

The matched statistic-transfer comparator uses the same binary tables, donors, splits, recipient margins, selection stage, and deviance-per-cell outcome. Stephenson pools donor statistics with the Paule–Mandel random-effects estimator [5]. For determinant $D_N = N_{00}N_{11} - N_{01}N_{10}$, the signed root Pearson statistic is

$$R_P = \text{sign}(D_N) \sqrt{X^2} = D_N \sqrt{\frac{n}{r_0 r_1 c_0 c_1}}. \quad (15)$$

The signed root likelihood-ratio deviance statistic is

$$R_G = \text{sign}(D_N) \left\{ 2 \sum_{i,j: N_{ij} > 0} N_{ij} \log(N_{ij}/E_{ij}) \right\}^{1/2}, \quad E_{ij} = r_i c_j / n. \quad (16)$$

Statistics are divided by $\sqrt{n}$ before donor pooling, multiplied by α , and restored to the recipient $\sqrt{n}$ scale before fixed-margin reconstruction. Stephenson used Paule–Mandel pooling with inverse-variance-plus-heterogeneity weights and selected the statistic, its raw or conditional-null-centered form, and α on the pilot. GSE239452 used equal donor means and selected a raw statistic and α by calibration cross-validation. COMBAT and BMMC used equal donor means and evaluated both centering rules. The hierarchical method transported conditional log odds and reconstructed the noncentral-hypergeometric expectation.

2.8 Fitted interaction baselines

For each entity e , the common-effect conditional maximum-likelihood estimate solves

$$\sum_d \mathbb{E}_{\theta_e} [A_{de} \mid m_{de}] = \sum_d A_{de}$$

The implementation returns the corresponding extended estimate at a support boundary and otherwise requires an absolute score below 10^{-8} . This estimator has no donor deviations, ridge penalty, or graph penalty.

The pooled-table baseline first sums the source-donor tables for each entity and computes the interaction coefficient from the saturated model

$$\log \mathbb{E}[N_{ij}] = \beta_0 + \beta_i^R + \beta_j^C + \beta^I \mathbb{I}\{i = 1, j = 1\}.$$

For a 2×2 table, β^I is the pooled sample log odds ratio, which differs from the stratified conditional estimate when donor margins vary. In the original post hoc comparison, both fitted coefficients select α on development donors and reconstruct held tables with Eq. (3). The pooled method is therefore pooled-table log odds with conditional reconstruction, not a Poisson plug-in prediction.

A separate post hoc comparator fixes the same pooled saturated-model interaction and profiles recipient row and column parameters. This gives the unique positive continuous table with the recipient margins and transported odds ratio, without Eq. (3). The original calibration and pilot donors select α before the interaction is refitted on all development donors.

2.9 Conditional random-effects prediction

For each marker pair, the classical random-effects model assumes $\theta_{de} \mid m_{de} \sim \mathcal{N}(\mu_e, \tau_e^2)$, with the same interaction distribution at every source margin, and estimates its parameters by conditional marginal likelihood [6]:

$$(\hat{\mu}_e, \hat{\tau}_e^2) = \arg \max_{\mu, \tau^2 \geq 0} \sum_d \log \int p(A_{de} \mid m_{de}, \theta) d\mathcal{N}(\mu, \tau^2)(\theta).$$

The common-effect model is the boundary $\tau_e^2 = 0$. Two predictions from the fitted random-effects distribution were compared: conditional expectation at $\alpha \hat{\mu}_e$ and the mean of conditional expectations over $\mathcal{N}(\alpha \hat{\mu}_e, \alpha^2 \hat{\tau}_e^2)$. The latter is a model for the interaction law conditional on recipient margins, not an unconditional law inferred by eliminating marginal nuisance parameters. Predicting from the distribution of a new unit's effect is standard in random-effects analysis [7].

The reanalysis retained the original 12-donor calibration, 24-donor pilot, and 56-donor recipient split. The pilot selected α separately for each prediction rule from $\{0.5, 0.75, 1, 1.25\}$; final fits used all 36 source donors. Predictions with $\alpha = 1$ were also retained. The hierarchy used its original penalties, and the common-ridge comparison used the same population penalty of 0.01. All comparisons are post hoc.

Predictive upper-left-count intervals were equal-tail intervals from the conditional or mixture law. Coverage and interval width were averaged over informative pairs within each recipient, then equally over recipients. The fitted source parameters were held fixed. Paired bootstrap intervals for loss comparisons used 20,000 physical-donor resamples.

2.10 Ground-truth estimator simulation

The factorial simulation crossed donor heterogeneity with recipient-margin shift. Each of 128 seeded replicates contained 12 source donors, 12 recipients, 16 independent binary entity tables per donor, and 128 observations per table. Counts followed the Fisher noncentral-hypergeometric distribution at known population log odds. Source donor log odds either equaled the population value or had independent normal deviations with standard deviation 1.2; recipient log odds equaled the population value without donor deviations. Recipient margins were source-like or shifted toward opposing row and column prevalences, using common random streams across margin conditions.

Across scenarios, the hierarchical estimator used heterogeneity penalty 0.1, ridge penalty 0.01, graph penalty zero, and $\alpha = 1$. The common-effect estimator fitted one conditional log odds across source donors with the same scaled population ridge penalty of 0.01; raw and conditional null-centered signed-root deviance statistics also used $\alpha = 1$. Recipient outcomes entered only the final deviance calculation. Intervals used 20,000 paired bootstrap draws over replicates.

2.11 Recipient heterogeneity

The extension crossed source and recipient log-odds standard deviations of 0 and 1.2 with both margin conditions. Each scenario used 128 replicates. The source draws and estimator settings were unchanged; all 256 source fits converged. Oracle predictions used the known population mean and recipient standard deviation, with Gauss–Hermite integration over recipient log odds.

2.12 Transfer-risk stress test

The stress test fixed the source field at its known value and crossed bounded mean-zero recipient drift, transport multipliers $\{0.5, 0.75, 1, 1.25\}$, drift standard deviations $\{0, 0.25, 0.5, 1\}$, and three informative margins. Finite-

Table 2: Ground-truth simulation: mean recipient multinomial deviance per cell. Parentheses give paired 95% bootstrap intervals; lower is better.

Donor effects	Recipient margins	Hierarchical	Common conditional	Raw statistic	Hierarchical comparison
Heterogeneous	source-like	0.014931	0.015212	0.015508	1.85% below common (1.52 to 2.18%); 3.72% below statistic transfer (3.28 to 4.15%)
Heterogeneous	shifted	0.011591	0.011714	0.033587	1.05% below common (0.84 to 1.26%); 65.49% below statistic transfer (63.02 to 67.72%)
Homogeneous	source-like	0.008752	0.008660	0.008688	1.06% above common (0.61 to 1.51%); 0.75% above statistic transfer (0.21 to 1.28%)
Homogeneous	shifted	0.008674	0.008640	0.029427	0.40% above common (0.13 to 0.67%); 70.52% below statistic transfer (69.21 to 71.74%)

Table 3: Recipient-heterogeneity extension. Relative reductions in deviance with paired-replicate 95% intervals; positive values favor the first method. Oracle comparisons use known parameters, not fitted estimates.

Source SD	Recipient SD	Margins	Hierarchy vs. common ridge	Oracle mixture vs. plug-in
0	0	source-like	-1.06% (-1.51, -0.61)	0%
0	0	shifted	-0.40% (-0.67, -0.13)	0%
0	1.2	source-like	0.31% (0.20, 0.42)	0.26% (0.14, 0.37)
0	1.2	shifted	0.28% (0.17, 0.39)	2.21% (1.89, 2.54)
1.2	0	source-like	1.85% (1.52, 2.19)	0%
1.2	0	shifted	1.05% (0.83, 1.26)	0%
1.2	1.2	source-like	0.41% (0.33, 0.49)	0.32% (0.22, 0.42)
1.2	1.2	shifted	0.39% (0.31, 0.47)	2.25% (1.88, 2.61)

support summation gave the expected conditional KL risk, numerical quadrature evaluated its curvature-integral form, and 10^6 seeded conditional draws per drift state gave an independent Monte Carlo estimate.

All 48/48 conditions satisfied the Fisher-information and finite-support envelopes. Exact KL and its curvature integral differed by at most 1.91×10^{-14} . Monte Carlo intervals covered the exact risk in 48/48 conditions, with maximum absolute $z = 1.833$. At $\alpha = 1$, increasing the drift standard deviation from 0 to 1 raised count-level expected log loss from 0 to 2.835 for balanced margins, 1.504 for row-sparse margins, and 1.322 for overlap-shifted margins; each table contained 96 observations.

3 Secondary multistate methods

3.1 Permutation-centered fields

For linked state vectors $(U_i, V_i)_{i=1}^n$, let $h_\alpha(U, V)$ denote Eq. (1) applied to the joint count table after adding pseudocount α . With B independent permutations of the V labels, the centered field is

$$\hat{\Theta} = h_\alpha(U, V) - B^{-1} \sum_{b=1}^B h_\alpha(U, V^{(b)}). \quad (17)$$

Under the fixed-margin link null, observed and permuted matchings are identically distributed conditional on both label multisets, so Eq. (17) has conditional mean zero. If S^2 is the coordinatewise sample variance of the B reference fields, $(1 + 1/B)S^2$ is an unbiased estimator of the variance of an observed-minus-reference field. Multistate analyses use $B = 64$, $\alpha = 0.5$, and one independent destroyed-link permutation.

For perturbation g and control 0 under challenge 1 and vehicle 0, the factorial field is $(\Theta_{g1} - \Theta_{g0}) - (\Theta_{01} - \Theta_{00})$. Independent arm-permutation variances are added before converting them to clipped relative precision,

$$w_{gj} = \text{clip} \left(\frac{1/v_{gj}}{\text{median}_{g,j}(1/v_{gj})}, 0.05, 20 \right).$$

3.2 Structured field estimator

For target-by-coordinate matrix Y , positive weight matrix W , positive-semidefinite hypergraph Laplacian L_H , and nonnegative penalties, the estimator minimizes

$$F_Y(M) = \frac{1}{2} \|Y - M\|_W^2 + \lambda_* \|M\|_* + \lambda_G \text{tr}(M^\top L_H M), \quad (18)$$

where $\|A\|_W^2 = \sum_{g,j} W_{gj} A_{gj}^2$. The normalized hypergraph Laplacian is [8]

$$L_H = I - D_v^{-1/2} H D_e^{-1} W_e H^\top D_v^{-1/2}.$$

Isolated vertices receive zero rows and columns and are unaffected by the graph penalty.

Each covered target formed an unweighted hyperedge with its five highest-cosine covered neighbors. The benchmark used frozen scGPT [9] and gene2vec [10] embeddings; label-permuted scGPT supplied the shuffled-graph control, and targets absent from an embedding were isolated. In PerturbSci, Frangieh, and Papalexi rotations, the other two units selected penalties by symmetric cross-prediction mean squared error. The nuclear penalty was a fraction of the largest training-matrix singular value, on grid $\{0, 0.01, 0.03, 0.1, 0.3, 0.6, 1\}$; the graph grid was $\{0, 0.05, 0.2, 1, 5, 20\}$. Nuclear-only, graph-only, and joint fits used the corresponding grids. The fixed variance-weighted fit used nuclear fraction 0.1 and graph penalty 5; scalar shrinkage used 26 multipliers from 0 to 1.25, and endpoint ridge used $\{10^{-4}, 10^{-3}, 10^{-2}, 0.1, 1, 10, 100\}$. MultiPerturb, PerturbFate, and Arce used the same five-neighbor graph and fixed penalties.

Proposition S5 (Uniqueness and nonexpansive stability). *If $\min_{g,j} W_{gj} > 0$, Eq. (18) has a unique minimizer. For two input matrices Y_1, Y_2 with the same weights, hypergraph, and penalties, their minimizers satisfy*

$$\|\widehat{M}_1 - \widehat{M}_2\|_W \leq \|Y_1 - Y_2\|_W. \quad (19)$$

Proof. The weighted quadratic loss is strongly convex and both penalties are convex, so their sum has a unique minimizer. Let R denote the combined penalty and choose $G_k \in \partial R(\widehat{M}_k)$. From $W \odot (\widehat{M}_k - Y_k) + G_k = 0$, subtracting the two optimality equations and taking the Frobenius inner product with $\widehat{M}_1 - \widehat{M}_2$ gives, by monotonicity of the convex subdifferential,

$$\|\widehat{M}_1 - \widehat{M}_2\|_W^2 \leq \langle W \odot (Y_1 - Y_2), \widehat{M}_1 - \widehat{M}_2 \rangle.$$

Weighted Cauchy-Schwarz proves Eq. (19). $\square$

Accelerated proximal gradient descent soft-thresholds singular values in the proximal step and handles the weighted residual and graph quadratic form in the smooth gradient. Iteration stops at a relative Frobenius step below 10^{-8} ; nonconverged estimates are rejected.

4 Public binary studies

Table 4: Binary RNA-protein studies.

Study	Accession	Linked object	Retained set	Unit and support rule
COMBAT	Zenodo 6120249	RNA-ADT states in one cell	81 ordered pairs	12 Oxford calibration and 24 Oxford pilot samples; development comparison only
BMMC	GSE194122	RNA-ADT states in one cell	100 ordered pairs	two fit donors and four bridge-donor batches; development comparison only
Stephenson	E-MTAB- 10026	RNA-ADT states in one cell	81 ordered pairs	12 Cambridge calibration, 24 Cambridge pilot, and 56 Newcastle held donors

Study	Accession	Linked object	Retained set	Unit and support rule
Oh pregnancy cohort	GSE239452	RNA-ADT states in one cell	81 ordered pairs	7 nonpregnant calibration, 8 nonpregnant pilot, and 9 pregnant held donors; post-access corrected replication
Babdor immunomicrobiome	GSE314416	RNA-ADT states in one cell	81 primary and 576 secondary ordered pairs	12 DB1 calibration and 20 pilot donors from DB1 and DB2; 77 donors from DB3 to DB7 remained unevaluated after the pilot criterion failed
Kelley biliary-cancer trial	GSE317605	RNA-ADT states in one cell	256 ordered pairs	7 calibration patients at four visits; 3 pilot and 8 held patients remained unevaluated after the calibration criterion failed

4.1 COMBAT and BMMC development panels

The COMBAT analysis used the official linked RNA-ADT object and sample manifest [11]. Metadata-only stratification assigned 12 Oxford samples to calibration and 24 to development. Each contributed 512 cells selected deterministically without using assay counts and 81 ordered pairs from nine matched markers.

The BMMC analysis used GSE194122 [12] as distributed with SCMMIB [13]. Donors 11466 and 19593 supplied the fit, and four batches from bridge donor 15078 supplied development evaluation. It included 4,607 cells and 100 ordered pairs from ten matched markers. Evaluation units in both panels participated in model selection.

4.2 Stephenson held-site confirmation

The Stephenson study profiled COVID-19 immune responses across UK sites [14]. Cambridge and Newcastle used BioLegend TotalSeq-C panel 99813; the protocol excluded 11 Sanger/London donors assayed with panel 99814 and six LPS-stimulated Newcastle donors. For each eligible donor, a deterministic rule independent of assay counts selected 512 cells. RNA state was raw UMI count above zero, and within-sample median splitting of each ADT marker fixed its margin at 256/256. Nine matched markers defined 81 ordered pairs, of which at least 64 had to be informative.

The prespecified split contained 12 Cambridge calibration donors (nine COVID-19 and three normal), 24 Cambridge pilot donors (18 COVID-19 and six normal), 11 unused Cambridge donors, and 56 Newcastle held donors. From the grid in Table 1, the pilot selected heterogeneity 0.1, ridge 0.01, graph 0, and $\alpha = 1$; the strongest statistic method was uncentered signed-root deviance with $\alpha = 0.75$. The final source fit used all 36 development donors.

Pilot mean deviance per cell was 0.005079 for the hierarchical estimator and 0.007761 for the selected statistic, a 34.55% reduction (paired donor-bootstrap loss-difference interval -0.003568 to -0.001827 ; 21/24 donors favorable). Relative to destroyed links, the reduction was 71.88% (-0.015650 to -0.009901 ; 23/24 favorable). Prediction then used Newcastle RNA and protein margins alone; paired Newcastle measurements entered once, during evaluation.

4.3 Matched predictive reanalysis

All 81 marker pairs yielded conditional random-effects fits in both the 12-donor calibration set and the full 36-donor source set. The pilot selected $\alpha = 0.75$ for both random-effects prediction rules and $\alpha = 1$ for the hierarchy and common-ridge model. Table 5 reports every fitted comparison, including the unscaled predictions.

The saved mean pilot losses for the mixture were 0.0071258792, 0.0053158800, 0.0053171639, and 0.0066445406 at $\alpha = 0.5$, 0.75, 1, and 1.25, respectively. The selected value improved pilot loss over $\alpha = 1$ by 1.284×10^{-6} , or 0.024%.

To assess multiplier-selection uncertainty, we recomputed donor-level losses at all four original multipliers from the saved source fits. Each of 20,000 bootstrap draws independently resampled 24 pilot donors and 56 recipients, selected the minimum-loss multiplier on its pilot resample, and scored the corresponding recipient predictions. The source fit was not re-estimated. Mixture and plug-in predictions used the same donor draws and selected their multipliers separately. This is a post hoc sensitivity to the calibration sample, not a source-refit bootstrap or an additional held cohort.

For the mixture, selection fractions at $\alpha = 0.5$, 0.75, 1, and 1.25 were 0.01%, 46.30%, 53.69%, and 0%. The plug-in fractions were 0.04%, 67.19%, 32.77%, and 0%. The central 95% bootstrap range for mixture improvement over the original hierarchy was -9.95% to 11.02% ; 46.305% of draws favored the mixture. The corresponding range for the mixture against the separately calibrated plug-in was -19.70% to 5.82% , with 79.07% of draws favorable.

These ranges include both multiplier reselection and recipient resampling; the following intervals condition on the originally selected multipliers.

The calibrated mixture reduced mean deviance by 8.46% relative to the original hierarchy (95% paired-donor interval, 5.42–11.55%), by 11.79% relative to common ridge (7.43–16.10%), and by 24.44% relative to the original signed-statistic comparator (16.03–31.16%). At the same fitted distribution and multiplier, integration reduced loss by 1.33% relative to the population plug-in (0.72–1.96%); 38/56 donors were favorable (one-sided sign-test $P = 0.00523$). The paired loss-difference interval was -0.000236 to -0.0000771 .

Table 5: Post hoc predictive comparison on the original 56 recipients and the 54-recipient selected-panel sensitivity. The latter removes two donors without changing fits, multipliers, or predictions. Loss is mean deviance per cell.

Estimator and prediction	α	56-donor loss	54-donor loss
Original hierarchy	1	0.012197	0.011393
Common conditional, ridge 0.01	1	0.012657	0.011713
Common conditional, unpenalized	1	0.013000	0.011965
Normal random effects, plug-in	0.75	0.011315	0.010723
Normal random effects, mixture	0.75	0.011165	0.010622
Normal random effects, plug-in	1	0.013865	0.012711
Normal random effects, mixture	1	0.013137	0.012098

After the selected-panel exclusion, the calibrated mixture improved on the original hierarchy by 6.77% (4.41–9.26%), common ridge by 9.31% (5.91–12.85%), and signed-statistic transfer by 28.86% (24.44–33.32%). The matched integration gain was 0.95% (0.55–1.36%), with 36/54 donors favorable ($P = 0.00992$).

Integration widened count intervals and increased coverage (Table 6). On the full recipient set, coverage rose from 70.01% to 83.17%, with mean interval width increasing from 15.74 to 22.90 cells. The 54-donor sensitivity gave 70.43% and 83.62% coverage, with widths of 15.66 and 22.69 cells.

Table 6: Coverage and width of nominal 95% predictive count intervals on 56 recipients. Parentheses give 95% donor-bootstrap intervals. Each recipient contributes equally after averaging over informative marker pairs.

Calibrated random-effects prediction	Coverage (%)	Width (cells)
Population plug-in	70.01 (67.04–73.04)	15.74 (15.38–16.09)
Predictive mixture	83.17 (80.72–85.60)	22.90 (22.02–23.74)

4.4 GSE239452 cohort replication

GSE239452 measured pregnancy-specific immunity after SARS-CoV-2 infection [15] and links GSE239450 gene expression with GSE239449 antibody-derived tags. Eligibility criteria excluded donor 100 for fewer than 512 shared cells and donor 78 for missing ADT data. The remaining split comprised seven nonpregnant calibration donors, eight nonpregnant pilot donors, and nine pregnant recipients: three controls, three asymptomatic infections, and three severe infections.

A deterministic rule independent of assay counts selected 512 shared barcodes per donor; state definitions, marker panel, pair set, and support threshold matched Stephenson. Seven-fold calibration selected heterogeneity 1, ridge 0.01, graph 0, and $\alpha = 1$; the comparator was uncentered signed-root deviance with $\alpha = 0.75$. After correction of its inverse reconstruction, pilot mean deviance was 0.008197 versus 0.010247, a 20.01% reduction (-0.004348 to -0.000546 ; 7/8 donors favorable). The destroyed-link reduction was 68.81% (-0.025268 to -0.012083 ; 8/8 favorable). The post-access correction changed only statistic-transfer predictions; primary and destroyed-link predictions and held tables were unchanged. GSE239452 is therefore a post-access corrected cohort replication.

The fixed-interaction Poisson analysis selected $\alpha = 1$ on 15 development donors. Its saturated source fit reconstructed all 81 pooled tables with maximum normalized cell error 1.78×10^{-16} .

4.5 GSE314416 held-pool attempt

GSE314416 profiled healthy participants across seven acquisition pools [16]. The prespecified split assigned 12 DB1 donors to calibration, 20 donors from DB1 and DB2 to the pilot, and 77 donors from DB3 to DB7 to held-pool evaluation. The primary 81-pair panel matched Stephenson; a secondary 24-marker panel contained 576 ordered pairs. Selection independent of assay counts retained 512 cells per donor.

Primary pilot loss was 0.001927, versus 0.001936 for selected signed Pearson-statistic transfer: a 0.49% reduction (paired donor interval -0.0000182 to -0.0000011 ; 13/20 favorable). The destroyed-link reduction was 0.45% (-0.0000390 to 0.0000223 ; 11/20 favorable), and the secondary panel reduced statistic-transfer loss by 2.54%. Neither panel reached the fixed 5% threshold, so the 77 held donors were not evaluated.

4.6 GSE317605 longitudinal calibration

GSE317605 profiled longitudinal PBMC RNA and surface proteins during a phase II trial of pembrolizumab plus granulocyte-macrophage colony-stimulating factor in advanced biliary cancer [17]. The split, fixed before assay-value access, assigned seven complete patients to calibration, three complete patients to the pilot, and eight patients to held evaluation. Five held patients had complete visit series and three had partial series. Each calibration visit retained 192 cells selected independently of assay counts and 256 ordered pairs from 16 matched RNA and ADT markers.

For visit v at time $t(v)$ and marker pair e , the longitudinal model decomposed the log odds into a time-specific population field $\beta_{t(v),e}$ and visit deviation u_{ve} . With V_t denoting visits at time t , it minimized

$$Q_{\text{long}} = \sum_t \frac{1}{|V_t|} \sum_{v \in V_t} \left\{ \sum_e \ell_{ve}(\beta_{te} + u_{ve}) + \frac{\eta}{2} \|u_v\|_2^2 \right\} + \frac{0.01}{2} \|\beta\|_F^2 \\ + \frac{\gamma}{2} \sum_t \beta_t^T L_{\text{pair}} \beta_t + \frac{\tau}{2} \sum_{t=1}^3 \|\beta_{t+1} - \beta_t\|_2^2, \quad (20)$$

where ℓ_{ve} is the exact fixed-margin negative log likelihood. The marker operator was $L_{\text{pair}} = L_H \otimes I + I \otimes L_H$, with L_H the normalized Laplacian of a fixed marker hypergraph; temporal smoothing followed the path from T01 to T04. The direct comparator fitted one time-specific interaction per pair by ridge-penalized profiled Poisson likelihood, then profiled row and column terms at each target margin. Both methods used identical patient-disjoint folds and transport grids. Loss was averaged across pairs within each visit, then across visits within each patient, with patients weighted equally.

Leave-one-patient-out calibration crossed donor-deviation penalties $\{0.1, 1, 10\}$, hypergraph and temporal penalties $\{0, 0.05, 0.2\}$, and transport multipliers $\{0, 0.5, 0.75, 1, 1.25\}$. The matched classical comparator crossed ridge penalties $\{0.001, 0.01, 0.1, 1\}$ with the same transport grid. The selected field fit used deviation penalty 0.1, temporal penalty 0.2, hypergraph penalty zero, and transport 0.75. The Poisson fit selected ridge penalty 1 and transport 0.75. Mean deviance was 0.006530, compared with 0.006609 for Poisson prediction. The reduction was 1.19%, with favorable results in 6/7 patients and 4/4 time points. Relative to destroyed links, deviance fell by 17.39%, with favorable results in 7/7 patients. Progression required at least 5% improvement over Poisson and strict improvement over a separately retuned graph-zero fit. The primary and graph-zero fits were identical, so the calibration criterion was not met and the pilot and held assay values were not accessed.

5 Binary conditional transfer results

5.1 Held outcomes

The Stephenson pilot required at least 5% lower donor-equal mean deviance than both statistic-transfer and destroyed-link controls, a 20,000-draw paired donor-bootstrap interval with upper endpoint below zero, and at least 19/24 favorable donors. Held-site confirmation further required at least 45/56 favorable donors and one-sided sign-test $P \leq 0.025$. GSE239452 used corresponding favorable-donor thresholds of 7/8 in pilot and 8/9 in held evaluation; its corrected statistic comparison is descriptive because repair followed outcome access.

Table 7: Held comparison with selected signed-statistic transfer.

Study and evidence role	Hierarchical	Statistic transfer	Reduction	Difference interval	Favorable donors
Stephenson, pre-outcome held site	0.01220	0.01478	17.46%	-0.00413 to -0.00080	50/56; sign $P = 5.09 \times 10^{-10}$
GSE239452, post-access corrected replication	0.00851	0.01413	39.81%	-0.00710 to -0.00423	9/9

Intervals are 20,000-draw paired physical-donor bootstrap intervals for hierarchical-minus-comparator mean deviance per cell.

Both studies selected $\lambda_G = 0$, making the graph-zero diagnostic identical to the selected fit.

5.2 Stephenson linkage and lineage sensitivity

The prespecified destroyed-link control deterministically permuted complete ADT profiles within each source donor. Of 32 independent post hoc repetitions, 29 met the numerical criterion, and every retained permutation had higher

Table 8: Held comparison with destroyed links.

Study and evidence role	Hierarchical	Destroyed	Reduction	Difference interval	Favorable donors
Stephenson, pre-outcome held site	0.01220	0.02274	46.36%	-0.01323 to -0.00749	51/56
GSE239452, post-access corrected replication	0.00851	0.04029	78.89%	-0.03567 to -0.02816	9/9

Intervals and loss definitions follow Table 7.

mean held loss than intact linkage. Against the donor-wise mean of retained controls, intact linkage reduced loss by 47.30% (paired donor interval 37.56% to 54.59%; loss-difference interval -0.01367 to -0.00795; 51/56 favorable, $P = 5.85 \times 10^{-11}$). Nested resampling over source permutations and held donors gave a loss-difference interval of -0.01371 to -0.00791.

At a common 64-cell donor budget, 54 T-cell, 52 B-cell, and 51 myeloid recipients were cell eligible. Pair-support filtering excluded 47, 42, and 5, leaving 7, 10, and 46 scorable donors, respectively. The fixed configurations reduced loss by -3.41% in T cells and -0.78% in B cells. The myeloid reduction was 15.51% (5.51% to 24.61%), with 21/46 favorable ($P = 0.769$). The site split and estimator settings remained fixed.

5.3 Composition-preserving and protein-threshold sensitivities

Eight additional source permutations shuffled intact nine-protein vectors within each donor and annotated cell type. All eight hierarchical fits converged. Their mean held-site loss was 0.011867, compared with 0.012197 for intact linkage. The relative reduction for intact linkage was -2.78% (paired-donor 95% interval, -9.05 to 2.87%); 31/56 donors favored intact linkage ($P = 0.252$, one-sided sign test). Each permutation had lower mean loss than intact linkage, with every paired loss-difference interval including zero. Joint resampling of source permutations and held donors gave a difference interval of -0.000336 to 0.001087.

Fixed source thresholds were raw ADT counts greater than 22, 4, 1, 3, 1, 9, 3, 30, and 46 for CD4, CD7, CD14, CD19, CD33, CD38, CD44, CD47, and CD52, respectively. These values were pooled source medians. All 81 pairs met the source-support criterion. Among 54 scorable recipients, hierarchical loss was 0.007927 and common conditional loss was 0.008147, a 2.70% reduction (95% interval, 0.72-4.73%). The loss-difference interval was -0.000396 to -0.0000567; 31/54 donors were favorable ($P = 0.170$).

5.4 Selected-panel signal sensitivity

The two recipients without informative threshold-based pairs were C-8914 (sample MH8919226) and C-8939 (MH8919227). Across 512 cells and the nine selected protein markers, their raw ADT totals were 3 and 0. C-8914 had one count each for CD4, CD14, and CD52, in three separate cells; the other selected markers were zero. All selected ADT counts in C-8939 were zero. The next-lowest selected-panel total across the other 90 source and recipient donors was 6,560. These totals concern the selected markers, not the full deposited ADT panel.

A post hoc sensitivity excluded these two donors without refitting models or changing any prediction. The original 56-donor confirmation was retained. On the remaining 54 donors, hierarchical deviance was 0.011393, compared with 0.014930 for signed-statistic transfer, 0.011965 for common conditional estimation, and 0.011644 for pooled-table conditional reconstruction. Corresponding reductions were 23.69% (95% interval, 17.86-28.99%), 4.77% (2.99-6.53%), and 2.15% (0.72-3.62%). Favorable-donor counts were 50, 44, and 36; one-sided sign-test values were 1.90×10^{-11} , 1.69×10^{-6} , and 0.00992.

5.5 Fitted interaction baselines

Post hoc fitted-interaction comparisons included a stratified common-effect estimate, pooled-table log odds with conditional reconstruction, and a fixed-interaction Poisson model that profiles recipient marginal parameters.

5.6 Matched development comparisons

Table 10 compares the conditional estimator with signed-statistic transfer in two adaptive development panels; reported units participated in selection, and intervals are descriptive.

The corresponding destroyed-link losses were 0.046395 in COMBAT and 0.013351 in BMMC, reductions of 75.09% and 18.73%. Both selected graph penalty zero; the best positive-graph configurations were 1.9% and 5.6% worse, respectively.

Table 9: Held comparisons with fitted interaction baselines.

Study	Comparator	Loss, hierarchy; comparator	Relative change (95% CI)	Loss-difference 95% CI	Favorable; exact P
Stephenson	exact common conditional	0.01220; 0.01300	6.18% lower (3.84% to 8.47%)	-0.00121 to -0.00046	46/56; 6.23×10^{-7}
Stephenson	pooled-table log odds, conditional reconstruction	0.01220; 0.01263	3.39% lower (1.42% to 5.37%)	-0.00075 to -0.00017	38/56; 0.00523
GSE239452	exact common conditional	0.00851; 0.00832	2.20% higher (1.34% to 3.03%)	0.00010 to 0.00028	0/9; 1.000
GSE239452	pooled-table log odds, conditional reconstruction	0.00851; 0.00996	14.58% lower (11.94% to 17.16%)	-0.00181 to -0.00115	9/9; 0.00195
GSE239452	fixed-interaction Poisson prediction	0.00851; 0.00998	14.79% lower (12.18% to 17.37%)	-0.00184 to -0.00117	9/9; 0.00195

Relative and loss-difference intervals use 20,000 paired donor-bootstrap draws. Differences are hierarchical-minus-comparator deviance per cell. Exact P values are one-sided binomial sign tests with zero differences counted as nonfavorable. All comparisons are post hoc and descriptive.

Table 10: Adaptive conditional-versus-statistic comparisons.

Panel	Hierarchical	Selected statistic transfer	Reduction	Difference interval	Favorable units
COMBAT, 24 Oxford pilot samples	0.011558	uncentered deviance, 0.015605	25.94%	-0.005967 to -0.002248	17/24 samples
BMBC, 4 bridge-donor batches	0.010851	centered deviance, 0.013148	17.47%	-0.005036 to -0.000385	4/4 batches

5.7 Public analysis sequence

Each analysis separated development selection, prediction from recipient margins, and evaluation (Table 11).

Table 11: Public analysis plans and study outcomes.

Study	Planned split	Evaluation status	Result or stopping reason
Stephenson	12 calibration, 24 pilot, 56 held-site donors	Pre-outcome held-site evaluation	Passed both prespecified comparators and donor-level criteria.
GSE239452	7 calibration, 8 pilot, 9 cohort-shift donors	Post-access corrected comparison	Corrected statistic inversion; reported as descriptive replication.
GSE314416	12 calibration, 20 pilot, 77 held-pool donors	Pilot evaluation	Missed the 5% and favorable-donor criteria; held pools were not evaluated.
GSE317605	7 calibration, 3 pilot, 8 held patients	Calibration evaluation	Missed the Poisson and graph-zero criteria; later stages were not evaluated.
GSE179221	Donor-disjoint source, development, and held sets	Eligibility assessment	Cognate feature-axis uniqueness failed before model fitting.
GSE342939	28 source and 6 held donors	Source evaluation	Every tested configuration failed in at least one source fold; held donors were not evaluated.

6 Secondary multistate results

6.1 Data and preprocessing

Table 12: Multistate studies and evaluation units.

Study	Accession	Linked object	Evaluation unit and support
PerturbSci-Kinetics [18]	GSE218566	old and new RNA in one cell	three sequence-distinct guide rotations; 85 targets
Frangieh [19]	scPerturb, Zenodo 13350497	RNA and surface protein in one cell	three guides per target and at least 20 cells per context; 179 targets

Study	Accession	Linked object	Evaluation unit and support
Papalexi [20]	scPerturb, Zenodo 13350497	RNA and surface protein in one cell	three deposited replicates and at least 20 cells per target; 24 targets
MultiPerturb [21]	GSE277747	RNA and ATAC in one nucleus	two sequence-distinct guides and at least 30 cells per guide; 35 targets
PerturbFate [22]	GSE291147	pre-existing and nascent RNA in one cell	four technical dates and at least 20 cells in every arm; 49 targets
Arce [23]	GSE278572	cognate RNA-ADT states in one cell	donor A development and donor B evaluation; 28 targets
ReSisTrace [24]	GSE223003	barcode-linked sister lineages	two biological replicates; 186 to 420 lineages per arm

PerturbSci used archived old- and new-RNA states from an encoder fitted to pooled non-targeting controls; guide sequences and controls were divided deterministically into three rotations. Frangieh and Papalexi used harmonized scPerturb RNA and protein objects [25], with separate principal-component and three-state k -means encoders fitted to control cells. MultiPerturb used archived RNA and ATAC states in reciprocal held-guide directions. PerturbFate fitted a three-state encoder to pooled non-targeting controls across 108,114 cells. After Souporecell exclusions, Arce included 100,086 cells and 112 cognate RNA-ADT features; donor A fixed the estimator and donor B supplied evaluation.

ReSisTrace used the 128 highest-variance pretreatment genes and an eight-component principal-component analysis fitted across pretreatment libraries. Three states were fitted to 507 control lineages. Barcode-defined lineages linked pre- and post-treatment states within each condition and replicate; treatment fields subtracted handling-control fields. Intervals used 2,000 bootstrap draws over linked lineages within each deposited replicate.

6.2 Predictive comparisons

Table 13 reports the structured estimator, strongest observed standardized-RMSE baseline, and destroyed-link control for each predictive multistate panel. PerturbSci passed guide replication but not the strongest observed RMSE comparison.

Table 13: Matched multistate comparisons.

Panel	Method	Pearson r (interval)	Standardized RMSE (interval)
PerturbSci	structured field	0.525 (0.483, 0.566)	0.844 (0.819, 0.868)
	marginal + structured residual	0.545 (0.507, 0.584)	0.831 (0.801, 0.862)
	destroyed link	0.074 (0.003, 0.144)	1.106 (1.071, 1.141)
Frangieh	structured field	0.040 (0.013, 0.066)	1.009 (1.002, 1.017)
	scalar shrinkage	0.007 (−0.026, 0.041)	0.999 (0.998, 1.001)
	destroyed link	0.092 (0.064, 0.122)	1.187 (1.164, 1.210)
Papalexi	structured field	−0.025 (−0.150, 0.096)	1.037 (1.006, 1.076)
	zero field	NA	1.000 (1.000, 1.000)
	destroyed link	−0.044 (−0.158, 0.058)	1.223 (1.139, 1.317)
MultiPerturb	structured field	−0.224 (−0.330, −0.124)	1.047 (1.005, 1.095)
	zero field	NA	1.000 (1.000, 1.000)
	destroyed link	−0.011 (−0.065, 0.059)	1.218 (1.145, 1.299)
PerturbFate	structured field	−0.092 (−0.152, −0.035)	1.023 (1.010, 1.035)
	linear cross-covariance	−0.011 (−0.122, 0.092)	1.000 (0.995, 1.004)
	destroyed link	−0.011 (−0.078, 0.056)	1.125 (1.089, 1.167)
Arce donor B	structured field	0.412 (0.331, 0.494)	1.093 (1.012, 1.185)
	linear cross-covariance	0.304 (0.195, 0.404)	0.993 (0.989, 0.996)
	destroyed link	0.170 (0.049, 0.299)	1.194 (1.081, 1.317)

Parentheses give 95% target-resampling intervals.

All eight ReSisTrace condition-by-replicate arms had fixed-margin Monte Carlo randomization $P = 1/65$ and Benjamini-Hochberg (BH) $q = 0.0154$. Each observed statistic and 64 disjoint test permutations was centered against

the same independent 64-permutation reference mean. No treatment contrast had a positive lineage-bootstrap lower limit.

Table 14: ReSisTrace treatment-minus-control replication.

Treatment	Replicate cosine	2,000-draw lineage-bootstrap 95% interval	Link P	BH q
Carboplatin	-0.674	-0.965, 0.729	0.877	0.877
Olaparib	0.223	-0.876, 0.925	0.354	0.531
NK	0.952	-0.801, 0.968	0.015	0.046

6.3 PerturbSci sensitivity and functional analysis

Table 15: PerturbSci performance in each held guide rotation.

Held rotation	Pooled Pearson r	Mean target cosine	Standardized RMSE
0	0.506	0.539	0.857
1	0.540	0.579	0.843
2	0.544	0.544	0.830

Table 16: PerturbSci sensitivity of pooled held-guide correlation.

Parameter	Value	Pooled r
States per assay	2	-0.121
	3	0.525
	4	0.640
Reference permutations	16	0.494
	32	0.508
	64	0.525
	128	0.523
Pseudocount	0.25	0.339
	0.50	0.525
	1.00	0.718

Module analysis required each target to enter the top field-norm quintile in every held truth and prediction, excluding targets that also passed after link destruction. No target qualified, so pathway and complex enrichment were not tested; XRN2 passed only in the destroyed control.

Mean top-five cosine-neighbor recovery was 0.064 (95% target-resampling interval 0.050 to 0.079) for the structured fit and 0.051 (0.041 to 0.063) after link destruction. The paired difference was 0.013 (-0.003 to 0.030 ; 10,000-draw target-label permutation $P = 0.283$). Reactome neighbor-edge enrichment failed its controls. CORUM edge enrichment exceeded destroyed links by 0.018 (0.008 to 0.029) and was nominal against random labels ($P = 0.0269$), but not after correction across two annotation families ($q = 0.0538$). No module, neighborhood, or individual edge met all criteria.

7 Assay resolution and recipient coexpression

This section gives the theory and retrospective analyses added after the original fixed-margin benchmark. All resolution comparisons hold the source interaction fixed.

7.1 Sharp coarse prediction bound

Let x, y range over finite assay-profile supports, let $S(x, y)$ be finite, and consider positive joint distributions

$$q(x, y) = \exp\{S(x, y) + a(x) + b(y)\}. \quad (21)$$

The binary queries $U = f(x)$ and $V = g(y)$ each have two nonempty fibers and known frequencies $m_U, m_V \in (0, 1)$. For every rectangle crossing their four fibers, define

$$d = S(x_0, y_0) + S(x_1, y_1) - S(x_0, y_1) - S(x_1, y_0),$$

where $f(x_0) = g(y_0) = 0$ and $f(x_1) = g(y_1) = 1$. Write $L = \min d$ and $H = \max d$. The closure of compatible coarse log odds is $[L, H]$. The following proof also gives the minimax prediction under these two frequency constraints. For the known margins m_U, m_V , write the compatible binary table as

$$P(t) = \begin{pmatrix} 1 - m_U - m_V + t & m_V - t \\ m_U - t & t \end{pmatrix}.$$

Let $t(\theta)$ be the unique interior cell probability giving log odds θ , and put $l = t(L)$ and $h = t(H)$.

Every fine distribution in Eq. (21) satisfies

$$q(x_0, y_0)q(x_1, y_1) = e^d q(x_0, y_1)q(x_1, y_0).$$

Summing over rectangles in the four coarse blocks expresses the coarse odds ratio as a positively weighted average of e^d . Its log therefore belongs to $[L, H]$. Concentrate the fine RNA masses $1 - m_U, m_U$ on an extremizing x_0, x_1 and protein masses $1 - m_V, m_V$ on y_0, y_1 . Their two-by-two reconstruction attains the corresponding endpoint. Positive fine distributions approximate these concentrated distributions while preserving the coarse margins. The positive fine-marginal domain is connected, and its entropic reconstruction is continuous, so the closure of attainable coarse tables is exactly $\{P(t) : l \leq t \leq h\}$.

For a fixed prediction, KL loss is convex in the true table, so its maximum on this segment occurs at an endpoint. Moving a predicted table along the segment increases its loss to one endpoint and decreases its loss to the other. There is a unique equalizing table P_* , which is a mixture of the endpoints. For $l < h$, define entropy as $\mathcal{H}(P) = -\sum P \log P$. The minimax log odds are

$$\theta_* = \frac{\mathcal{H}(P(l)) - \mathcal{H}(P(h))}{h - l}, \quad P_* = P(t(\theta_*)).$$

For $l = h$, set $P_* = P(l)$.

The equalizing construction is an information-radius minimax rule [26]. To prove optimality over all predictions, including those with different margins, choose the weight w such that $P_* = wP(l) + (1 - w)P(h)$. For any prediction R ,

$$\begin{aligned} w \text{KL}(P(l) \| R) + (1 - w) \text{KL}(P(h) \| R) \\ = R_* + \text{KL}(P_* \| R), \end{aligned}$$

where $R_* = \text{KL}(P(l) \| P_*) = \text{KL}(P(h) \| P_*)$. The worst endpoint loss cannot be smaller than this weighted average. Equality is attained uniquely by P_* . All compatible populations induce the same independent unpaired Bernoulli observation laws with frequencies m_U, m_V . Convexity in the predicted distribution makes randomization unhelpful, so R_* is also a lower bound for every sample-based predictor observing only these unpaired binary states.

For the three-marker example in the main text, take the RNA distribution $r(z_1, z_2) = (1 + \rho z_1 z_2)/4$, with $|\rho| < 1$. The joint distribution

$$q(z_1, z_2, w) = r(z_1, z_2) \frac{e^{J z_2 w}}{2 \cosh J}$$

has all three individual means zero and the required fine interaction. Conditional expectation gives $\mathbb{E}(W | Z_1, Z_2) = Z_2 \tanh J$ and hence $\mathbb{E}(Z_1 W) = \rho \tanh J$. All probabilities are positive, so the example does not depend on structural zeros.

7.2 Response to a transferred interaction

The expansion below specializes the fixed-marginal response identities of Sei and Yano [27]. Fix positive marginal masses $r(x)$ and $c(y)$ on finite supports. Write $\bar{x} = x - \mathbb{E}_r X$, $\bar{y} = y - \mathbb{E}_c Y$, and $h = \bar{x}^\top B \bar{y}$. Centering changes the interaction only by additive assay functions. The reconstruction can therefore be written

$$q_t(x, y) = r(x)c(y) \exp\{th(x, y) + \alpha_t(x) + \beta_t(y)\}.$$

At $t = 0$, the adjusting functions vanish. After fixing their additive gauge, the finite marginal equations have a nonsingular linearization at independence, so their solution is analytic near zero. Differentiating the marginal constraints once gives zero conditional means for $\partial_t \log q_t$. Centering h in both arguments therefore gives $\partial_t \log q_t|_{t=0} = h$. Differentiating twice gives

$$\begin{aligned} \left. \frac{\partial_t^2 q_t}{rc} \right|_{t=0} &= h^2 - \mathbb{E}_c(h^2 | x) \\ &\quad - \mathbb{E}_r(h^2 | y) + \mathbb{E}_{rc} h^2. \end{aligned}$$

When multiplied by $\bar{x}_i\bar{y}_j$ and averaged under rc , the last three terms vanish. Define the third central moment tensors by $T_{iac}^x = \mathbb{E}_r(\bar{X}_i\bar{X}_a\bar{X}_c)$ and $T_{jbd}^y = \mathbb{E}_c(\bar{Y}_j\bar{Y}_b\bar{Y}_d)$. It follows that

$$\begin{aligned}\text{Cov}_{q_t}(X_i, Y_j) &= t(\Sigma_x B \Sigma_y)_{ij} \\ &+ \frac{t^2}{2} \sum_{a,b,c,d} B_{ab} B_{cd} T_{iac}^x T_{jbd}^y + O(t^3).\end{aligned}$$

For an explicit example, take three spins in each assay and let

$$r(x) = \frac{1 + \rho_x x_1 x_2 x_3}{8}, \quad c(y) = \frac{1 + \rho_y y_1 y_2 y_3}{8},$$

with $x_i, y_i \in \{-1, 1\}$ and $|\rho_x|, |\rho_y| < 1$. Every individual mean is zero and each within-assay covariance matrix is the identity. For diagonal $B = \text{diag}(b_1, b_2, b_3)$,

$$\mathbb{E}_{q_t}(X_1 Y_1) = t b_1 + t^2 \rho_x \rho_y b_2 b_3 + O(t^3).$$

The pairwise maximum-entropy replacement is uniform in each assay and loses the second term. Positive distributions with identical means and covariances can therefore give different RNA–protein associations under the same bilinear interaction.

7.2.1 Query discrepancy from covariance mismatch

Consider recipient and pooled reconstructions with the same binary marker frequencies $m, n \in (0, 1)$ and interaction $tx^\top B y$. Let k_R and k_P be their respective queried entries of $\Sigma_x B \Sigma_y$. Both query tables equal the independent table P_0 at $t = 0$, and

$$P_R(t) - P_P(t) = t(k_R - k_P)V + O(t^2),$$

where $V = \begin{pmatrix} 1 & -1 \\ -1 & 1 \end{pmatrix}$. Taylor expansion of KL divergence about the positive table P_0 gives

$$\begin{aligned}2 \text{KL}(P_R(t) \| P_P(t)) &= t^2 (k_R - k_P)^2 \sum_{u,v} \frac{V_{uv}^2}{(P_0)_{uv}} + O(t^3) \\ &= \frac{t^2 (k_R - k_P)^2}{m(1-m)n(1-n)} + O(t^3).\end{aligned}$$

The coefficient is the response score used to rank queries. If the recipient model describes the true joint law, it is also the leading excess deviance from substituting pooled coexpression. The empirical comparison tests the ranking at the fitted interaction strength without imposing that equality.

7.3 Matched resolution and covariance analyses

The three recipient cohorts contain 24 Cambridge donors, 56 Newcastle donors, and 15 T-ALL patients. The same nine RNA and nine protein markers give 81 ordered queries. All analyses retain the original deterministic 512-cell sample, donor-wide protein ranks, 256 adaptation cells, 256 scoring cells, and 12-donor source interaction with pairing-contrast penalty 10^{-4} . The paired source fit is the only source of cross-assay interaction information. Recipient adaptation uses separate-assay patterns and the scoring sample’s individual marker frequencies, not its pairings.

The means-only reconstruction constrains all 18 individual frequencies. Each one-pattern arm additionally constrains that assay’s full binary-profile histogram. The both-pattern arm constrains both histograms. Histograms receive four source-prior pseudocells and exponential tilts to the scoring-cell marker frequencies. Constant markers are restricted to their observed states. Matrix scaling fits complete pattern constraints, and moment matching fits individual-frequency constraints. These are reconstructions from the same interaction, not separately trained interaction models.

The covariance arm replaces each smoothed and tilted histogram with its maximum-entropy distribution subject to the same first and second moments. For binary profiles its log density contains unary and pairwise terms. This replacement can induce higher-order correlations; it fixes pairwise moments rather than forcing all higher-order moments to vanish. The retained-gain fraction is

$$\frac{L_{\text{means}} - L_{\text{covariance}}}{L_{\text{means}} - L_{\text{patterns}}},$$

where each L is the cohort’s mean donor deviance. It is a descriptive loss decomposition, not an estimate of a mediated biological effect.

7.4 Recipient-specific and pooled coexpression

The source pool uses the 12 source donors’ adaptation halves: 3,072 cells per assay. For each query person, the cohort pool uses every other recipient’s adaptation half and excludes the query person. It contains 5,888 cells in Cambridge, 14,080 in Newcastle, and 3,584 in T-ALL. Both controls use four prior pseudocells per contributing person, preserving the prior fraction, then undergo the same tilting and maximum-entropy reconstruction at the query person’s marker frequencies. The cohort pool uses recipient-cohort data but no scoring-cell patterns. These controls vary the origin of covariance without changing the interaction, marker frequencies, or scoring tables.

7.5 Query ranking and uncertainty

For RNA marker i and protein marker j , define

$$s_{ij} = \frac{\{(\Sigma_x B \Sigma_y)_{ij}^{\text{recipient}} - (\Sigma_x B \Sigma_y)_{ij}^{\text{pool}}\}^2}{m_i(1 - m_i)n_j(1 - n_j)}.$$

The source-strength comparator is $B_{ij}^2 m_i(1 - m_i)n_j(1 - n_j)$. Scores are zero for constant-marker queries. Both rankings use the fixed marker order to break ties. Pair-specific gain is the cohort-pool deviance minus recipient-covariance deviance, using unchanged predictions. The top 27, bottom 27, and source-ranked top 27 were specified before computing the pair-specific losses. Cumulative gain curves sum gains over people and selected queries, then divide by the sum over all 81 queries; they do not average person-specific percentages.

All loss endpoints average twice KL over every marker pair, including constants, then equally over people. Each comparison uses 20,000 paired person-level bootstrap draws. The four analysis families use seeds 906091–906094 and contain, respectively, nine pattern-resolution contrasts, six covariance-replacement contrasts, six recipient-versus-pool contrasts, and six query-ranking contrasts. For a family of M comparisons, the Bonferroni percentile endpoints are $0.025/M$ and $1 - 0.025/M$. These intervals condition on the fixed source fit and fitted pools.

Table 17: Query-ranking comparisons. Values are mean deviance-gain differences per selected pair. Intervals resample people and correct across six contrasts.

Cohort	Comparison	Difference	Favorable people	Bonferroni interval
Cambridge	Response top vs bottom third	0.001711	15/24	−0.000126 to 0.003883
Newcastle	Response top vs bottom third	0.004983	44/56	0.002657 to 0.007787
T-ALL	Response top vs bottom third	0.024259	14/15	0.014180 to 0.037793
Cambridge	Response vs source-ranked top third	0.000038	11/24	−0.001043 to 0.001154
Newcastle	Response vs source-ranked top third	0.001642	38/56	0.000612 to 0.002920
T-ALL	Response vs source-ranked top third	0.013881	14/15	0.007619 to 0.021405

7.6 Relation to full-profile interaction fitting

The earlier T-ALL comparison used all 15 patients, the same permitted recipient information, and additional protein-abundance coordinates. A common recipient adapter compared the no-prior CHAMPOLLION inverse-transport objective [28] with donor-wise joint fitting, query-targeted fitting, and a multinomial comparator. Mean losses were 0.008441, 0.008476, 0.008602, and 0.010630, respectively. Query-targeted fitting improved on the multinomial comparator by 19.07% (95% interval, 13.05–24.36%); its relative improvement over the CHAMPOLLION objective was −1.91% (−6.26 to 2.25%). This comparison did not meet its all-control superiority criterion. It neither benchmarks the complete CHAMPOLLION workflow nor establishes a better interaction learner. The subsequent binary resolution analyses hold one fitted interaction constant.

7.7 Computational verification

Separate implementations reconstructed all 380 initial resolution predictions, 190 covariance-preserving assay distributions and their 95 joint predictions, and 190 pooled-control predictions. Maximum absolute differences in predicted probabilities were 5.04×10^{-11} , 1.14×10^{-10} , and 1.27×10^{-10} , respectively. Verification recounted the scoring tables, checked exact marker margins, and checked source membership and exclusion of the query person from each cohort pool. Independently calculated query scores differed by at most 1.47×10^{-10} ; pair-specific gain and bootstrap-summary differences were below 4×10^{-16} . Fourteen tests cover marginal constraints, reconstruction equivalence, covariance replacement, pool construction, and query scoring.

Analysis plans preceded these computations but followed examination of earlier outcomes; the analyses are retrospective. The original public v2.0.4 benchmark is a separate release and does not contain the subsequent resolution, covariance, pooled-control, or query-ranking analyses.

8 Limitations

The 81 within-donor pair tables reuse cells and marker vectors, but the composite objective does not model their joint covariance. Held-donor bootstraps condition on the fitted source field, selected configuration, and deterministic cell subsamples. The original within-donor median ranks fixed every Stephenson protein margin at 256/256, including two recipients with almost no counts across the selected protein panel. The source-fixed threshold analysis permits protein-margin variation but does not establish biologically validated positivity thresholds. Both this analysis and the selected-panel exclusion followed outcome access.

Each primary destroyed-link comparison uses one deterministic within-donor permutation; the repeated-permutation sensitivity retained 29 of 32 fits, and three did not meet the convergence criterion. The eight within-cell-type permutations address a different null, preserving broad cell composition. Neither predictive interval family attained nominal 95% coverage. These intervals treat fitted interaction distributions as known and assume their transportability; they omit source-parameter uncertainty. The simulations use 16 independent entity tables and deliberately strong margin shift. Recipient interactions equal the population parameter in the original four scenarios and vary around it in the extension. Neither simulation defines a multivariate cell-level generative law.

The fitted normal model assumes a margin-independent conditional interaction law. Corollary S1 permits margin-dependent laws, but does not identify their recipient values from marginal counts. Its bound concerns mean-table regret, not absolute error or predictive interval coverage. The calibration sensitivity reselects the nearly tied multipliers but still conditions on the source fit; it does not account for source-estimation uncertainty. The composition-preserving control evaluated the hierarchy, not the random-effects mixture. Separately predicted pairwise tables need not be compatible with a single multivariate cell distribution.

Conditional likelihood removes row and column nuisance parameters but can lose finite-sample information about the odds ratio when one or both margins are random under the sampling design. The lost fraction tends to zero under standard large-sample conditions [29]. We did not estimate its magnitude in the sparse marker tables used here.

The original GSE239452 statistic inversion mishandled negative values near the lower feasible endpoint, producing 80 coordinate mismatches among 729 reconstructed tables. The corrected inverse produced none at tolerance 10^{-8} and maximum absolute coordinate error below 4.0×10^{-12} . It recomputed only statistic-transfer predictions; primary and destroyed-link predictions and held tables were unchanged. Because repair followed outcome access, GSE239452 is a post-access corrected replication, not a second pre-outcome confirmation.

9 Implementation and reproducibility

The v2.0.4 public release [30] contains the MIT-licensed estimator, analysis protocols, cohort and simulation results, study records, and method grids. Its verification tests reproduce the reported tables and figures from the derived results. Raw matrices remain available from the accessions in Tables 4 and 12 under their source licenses and are not redistributed.

The reanalysis uses the original 92 donor samples, 512 selected cells per donor, and nine markers per modality. Recovering their selected counts from the public source reproduced all 56 primary held-donor losses to a maximum absolute difference of 5.7×10^{-15} . Released aggregate contingency tables reproduce the biological sensitivities without downloading cell-level data. Assay-quality diagnostics and the unchanged original prediction losses reproduce the selected-panel exclusion. The analysis specification and computational provenance distinguish these post hoc analyses from the original held-site protocol.

The recipient-heterogeneity simulation reproduced the original zero-variance losses within 7.0×10^{-18} ; 40- and 80-node quadrature means agreed within 8.7×10^{-13} . Tests check predictive mixtures against independent adaptive integration, degenerate margins, zero variance, convergence failure, and the distribution-to-table regret bound.

The classical random-effects implementation uses `metafor` 5.0.1's exact conditional model, `CM.EL`, with no count correction. Tables with singleton support contribute likelihood one and are omitted from fitting. Fits were checked at two integration tolerances and against independent finite-support likelihood calculations. Identical informative donor tables admit an analytic zero-variance solution. Other near-boundary fits were checked with additional optimizers, the one-sided variance score, and a fixed variance profile. Predictive intervals use the fitted count law, not ancillary parameter standard errors. The common-ridge calibration fit solved the score equations to maximum absolute gradient 9.2×10^{-11} . The release records optimizer settings and fit provenance.

References

- [1] Liao JG. A hierarchical Bayesian model for combining multiple 2×2 tables using conditional likelihoods. *Biometrics*. 1999;55(1):268–272. <https://doi.org/10.1111/j.0006-341x.1999.00268.x>.
- [2] Coull BA, Agresti A. Generalized log-linear models with random effects, with application to smoothing contingency tables. *Statistical Modelling*. 2003;3(4):251–271. <https://doi.org/10.1191/1471082X03st059oa>.

- [3] Stijnen T, Hamza TH, Özdemir P. Random effects meta-analysis of event outcome in the framework of the generalized linear mixed model with applications in sparse data. *Statistics in Medicine*. 2010;29(29):3046–3067. <https://doi.org/10.1002/sim.4040>.
- [4] Kou SG, Ying Z. Asymptotics for a 2×2 Table with Fixed Margins. *Statistica Sinica*. 1996;6(4):809–829.
- [5] Paule RC, Mandel J. Consensus values and weighting factors. *Journal of Research of the National Bureau of Standards*. 1982;87(5):377–385. <https://doi.org/10.6028/jres.087.022>.
- [6] Jackson D, Law M, Stijnen T, Viechtbauer W, White IR. A comparison of seven random-effects models for meta-analyses that estimate the summary odds ratio. *Statistics in Medicine*. 2018;37(7):1059–1085. <https://doi.org/10.1002/sim.7588>.
- [7] Higgins JPT, Thompson SG, Spiegelhalter DJ. A re-evaluation of random-effects meta-analysis. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*. 2009;172(1):137–159. <https://doi.org/10.1111/j.1467-985X.2008.00552.x>.
- [8] Zhou D, Huang J, Schölkopf B. Learning with hypergraphs: clustering, classification, and embedding. In: *Advances in Neural Information Processing Systems 19*; 2006. p. 1601–1608.
- [9] Cui H, Wang C, Maan H, Pang K, Luo F, Wang B. scGPT: toward building a foundation model for single-cell multi-omics using generative AI. *Nature Methods*. 2024;21:1470–1480. <https://doi.org/10.1038/s41592-024-02201-0>.
- [10] Du J, Jia P, Dai Y, Tao C, Zhao Z, Zhi D. Gene2vec: distributed representation of genes based on co-expression. *BMC Genomics*. 2019;20(Suppl 1):82. <https://doi.org/10.1186/s12864-018-5370-x>.
- [11] COVID-19 Multi-omics Blood Atlas (COMBAT) Consortium. A blood atlas of COVID-19 defines hallmarks of disease severity and specificity. *Cell*. 2022;185(5):916–938.e58. <https://doi.org/10.1016/j.cell.2022.01.012>.
- [12] Luecken MD, Burkhardt DB, Cannoodt R, Lance C, Agrawal A, Aliee H, et al. A sandbox for prediction and integration of DNA, RNA, and proteins in single cells. In: *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks*. vol. 1; 2021. p. 1–13. Available from: <https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/158f3069a435b314a80bdc024f8e422-Abstract-round2.html>.
- [13] Fu S, Wang S, Si D, Li G, Gao Y, Liu Q. Benchmarking single-cell multi-modal data integrations. *bioRxiv*. 2025. Preprint. <https://doi.org/10.1101/2025.04.01.646578>.
- [14] Stephenson E, Reynolds G, Botting RA, Calero-Nieto FJ, Morgan MD, Tuong ZK, et al. Single-cell multi-omics analysis of the immune response in COVID-19. *Nature Medicine*. 2021;27:904–916. <https://doi.org/10.1038/s41591-021-01329-2>.
- [15] Oh DS, Kim E, Normand R, Lu G, Shook LL, Lyall A, et al. SARS-CoV-2 infection elucidates features of pregnancy-specific immunity. *Cell Reports*. 2024;43(11):114933. <https://doi.org/10.1016/j.celrep.2024.114933>.
- [16] Babdor J, Patel RK, Davidson B, Koser K, Noecker C, Rahim MK, et al. Immune-microbiome coordination defines interferon setpoints in healthy humans. *Cell*. 2026;189(10):3144–3163.e20. <https://doi.org/10.1016/j.cell.2026.02.003>.
- [17] Kelley RK, Bracci P, Gordan JD, Behr SC, Quandt Z, Atreya CE, et al. A Phase II Trial of Pembrolizumab plus Granulocyte-Macrophage Colony-Stimulating Factor in Advanced Biliary Cancer. *Clinical Cancer Research*. 2026. Published online 14 August 2026. <https://doi.org/10.1158/1078-0432.CCR-25-3987>.
- [18] Xu Z, Sziraki A, Lee J, Zhou W, Cao J. Dissecting key regulators of transcriptome kinetics through scalable single-cell RNA profiling of pooled CRISPR screens. *Nature Biotechnology*. 2024;42:1218–1223. <https://doi.org/10.1038/s41587-023-01948-9>.
- [19] Frangieh CJ, Melms JC, Thakore PI, Geiger-Schuller KR, Ho P, Luoma AM, et al. Multimodal pooled Perturb-CITE-seq screens in patient models define mechanisms of cancer immune evasion. *Nature Genetics*. 2021;53:332–341. <https://doi.org/10.1038/s41588-021-00779-1>.
- [20] Papalexi E, Mimitou EP, Butler AW, Foster S, Bracken B, Mauck WM, et al. Characterizing the molecular regulation of inhibitory immune checkpoints with multimodal single-cell screens. *Nature Genetics*. 2021;53:322–331. <https://doi.org/10.1038/s41588-021-00778-2>.
- [21] Yan RE, Corman A, Katgara L, Wang X, Xue X, Gajic ZZ, et al. Pooled CRISPR screens with joint single-nucleus chromatin accessibility and transcriptome profiling. *Nature Biotechnology*. 2025;43:1628–1634. <https://doi.org/10.1038/s41587-024-02475-x>.
- [22] Xu Z, Lu Z, Ugurbil A, Abdulraouf A, Liao A, Zhang J, et al. Mapping convergent regulators of melanoma drug resistance by PerturbFate. *Nature*. 2026;654:261–271. <https://doi.org/10.1038/s41586-026-10367-0>.
- [23] Arce MM, Umhoefer JM, Arang N, Kasinathan S, Freimer JW, Steinhart Z, et al. Central control of dynamic gene circuits governs T cell rest and activation. *Nature*. 2025;637:930–939. <https://doi.org/10.1038/s41586-024-08314-y>.
- [24] Dai J, Zheng S, Falco MM, Bao J, Eriksson J, Pikkusaari S, et al. Tracing back primed resistance in cancer via sister cells. *Nature Communications*. 2024;15:1158. <https://doi.org/10.1038/s41467-024-45478-7>.
- [25] Peidli S, Green TD, Shen C, Gross T, Min KH, Garda S, et al. scPerturb: harmonized single-cell perturbation data. *Nature Methods*. 2024;21:531–540. <https://doi.org/10.1038/s41592-023-02144-y>.
- [26] van Erven T, Harremoës P. Rényi divergence and Kullback–Leibler divergence. *arXiv preprint*. 2014. ArXiv:1206.2459.
- [27] Sei T, Yano K. Minimum information dependence modeling. *Bernoulli*. 2024;30:2623–2643. <https://doi.org/10.3150/23-BEJ1687>.

- [28] Samaran J, Peyré G, Cantini L. CHAMPOLLION: Robust Multi-Omics Integration via Inverse Optimal Transport Using Paired Cells. bioRxiv. 2026. <https://doi.org/10.64898/2026.04.28.721317>.
- [29] Schervish MJ, Kadane JB, Seidenfeld T, Stern RB. Information lost by conditioning on the random margins in a two-by-two table. Biometrika. 2025;112(4):asaf064. <https://doi.org/10.1093/biomet/asaf064>.
- [30] sushaan-k. Coupling Fields: public benchmark, version 2.0.4 [Internet]. GitHub; 2026. Available from: <https://github.com/sushaan-k/coupling-fields-benchmark/releases/tag/coupling-fields-v2.0.4-public-benchmark>. Accessed 5 September 2026.